

# Lab 6: Synchronizing Threads

Master M1 MOSIG – Université Grenoble Alpes & Grenoble INP

2018

## 1 Instructions

**This lab is not graded. The problems are to be solved on machine.**

Before starting working on this lab, you should have finished solving all exercises from the practical section of Lab 5.

## 2 Debugging multi-threaded applications

We have introduced the debugging tool `gdb` during the first lab. `gdb` allows step-by-step execution during which the user can explore the state of the memory. `gdb` can be useful to debug multi-threaded applications.

The file `prod_cons_bug.c` is supposed to implement a producer-consumer synchronization between threads. Unfortunately this code is buggy.

**Question 2.1:** *Run the program described in `prod_cons_bug.c`. What is the problem that you observe?*

The goal of this exercise is to use `gdb` to identify the bug. You can find a description of how to use `gdb` with threads here: [https://ftp.gnu.org/old-gnu/Manuals/gdb/html\\_node/gdb\\_24.html](https://ftp.gnu.org/old-gnu/Manuals/gdb/html_node/gdb_24.html). We summarize the main commands:

- `info threads`: Display a summary of all threads currently in your program.
- `thread threadno`: Make thread number `threadno` the current thread

Note also that you can interrupt an application running in `gdb` by sending it a `SIGINT` signal using `Ctrl-c`.

**Question 2.2:** *Fix the bug in the program using what you observe with `gdb`.*

### 3 Spinlock

In this exercise, you are asked to implement a spinlock to provide mutual exclusion for an *a priori* unknown number of threads.

**Question 3.3:** *What operations should be supported by the hardware to be able to solve the problem?*

**Question 3.4:** *Complete the provided file `spinlock.c` by implementing a spinlock using `Test_And_Set()`.*

**Note:** You can find the signature of the built-in atomic memory access operations supported by gcc at <https://gcc.gnu.org/onlinedocs/gcc-4.1.0/gcc/Atomic-Builtins.html>.

**Question 3.5:** *Implement a new spinlock, using this time `Compare_And_Swap()`.*

The locks you have implemented until now do not ensure fairness. It is rather simple to implement a *fair* spinlock using `Fetch_And_Add()`. Such a lock is actually called a *ticket* lock.

**Question 3.6:** *Implement a fair spinlock using `Fetch_And_Add()`. By fair, we mean here that if thread A calls `lock()` before thread B, then thread A should be granted access to the critical section before thread B.*

**Question 3.7:** *Design a simple test to confirm that the new lock is fair and that the previous versions were not.*

Running the default test included in file `spinlock.c` with a adequate number of threads should show you that the lock based on `Fetch_And_Add()` is much less efficient than the *unfair* spinlocks.

**Question 3.8:** *What is the number of threads at which the performance of the ticket lock starts to really degrade?*

**Question 3.9:** *Try to explain this behavior.*